

Blackletter — BMAD V4 Implementation Guide (One-Stop)

Project: Blackletter — GDPR Processor-Obligations Checker

Audience: Vibe CEO, PM, Architect, UX, PO, SM, Dev, QA

Scope: A single, step-by-step guide that stitches together all artifacts produced so far (Analyst → UX → Architect → PO → SM) and shows exactly how to execute the IDE stage with agents until MVP ships.

0) TL;DR Quickstart (15 minutes)

1. **Clone & run** the scaffold (Windows):

```
scripts\ps\dev.ps1
```

2. **Seed config** (already provided): `core-config.yaml` (LLM off, caps, OCR off).
3. **Shard docs** (creates per-section files for agents):

```
*shard-doc docs/prd.md prd
*shard-doc docs/architecture.md architecture
```

4. **Start new chat tabs** per agent in your IDE.
 5. **Scrum Master** → open **SM Story Pack v1** and choose **Story 1.1**.
 6. **Dev** builds 1.1, marks ready; **QA** verifies; loop stories **1.1 → 1.2 → 1.3 → 2.1 → 2.2 → 3.1 → 3.2 → 2.4 → 4.1 → 4.2 → 5.1 → 5.2**.
 7. **Export** a report from Findings; review Metrics wall; done.
-

1) What is a Story, and How the IDE Stage Works

- A **story** is a **testable feature slice** with: ID, epic, title, acceptance criteria, tasks, and tests.
 - **IDE stage** = planning docs become **code** through an agent loop: 1) **SM**: turns approved PRD items into **developer-ready packets** (already prepared in the Story Packs).
 - 2) **Dev**: implements a single story using those packets + dev-load-always files.
 - 3) **QA**: verifies acceptance criteria, runs tests, may refactor.
 - 4) Repeat until epics complete.
 - Always run each agent in a **fresh chat/tab** to keep context clean.
-

2) Artifact Map (what you already have)

- **Analyst (Mary):** *Hand-Off Pack v1* — acceptance wordings, weak-language lexicon seed, rulepack skeleton, KPIs, gold-set plan.
- **PM (John):** *PRD v1* — vision, epics, draft stories, KPIs, scope.
- **UX (Sally):** *UI/UX Spec v1* — IA, flows, tokens, components, wireframes, v0/Lovable prompts.

- **Architect (Winston):** *Architecture v1* — stack pins, source tree, API, detection engine, Windows commands.
- **PO (Sarah):** *PO Validation & Approvals v1* — checklist pass, API/schema freeze, story approvals, sharding plan.
- **SM (Sam):** *SM Story Pack v1* — dev-ready packets for 1.1, 1.2, 2.1, 2.2, 3.1.
- **Gap Closure Pack (Agents Complete v1):** remaining stories (1.3, 2.4, 3.2, 4.1, 4.2, 5.1, 5.2), full rulepack a-h, schema validation, gold set + scorer, core config, error codes, PII redaction, PS scripts, CI, QA checklists.

Keep these under `docs/` and referenced in IDE chats.

3) Repository Layout (from Architecture v1)

```
blackletter/
  apps/
    web/                # Next.js 14 app
    api/                # FastAPI service
  docs/
    prd.md
    architecture.md
    prd/                # (after sharding)
    architecture/       # (after sharding)
    stories/
    core-config.yaml    # authoritative config (LLM off, caps, OCR
off)
    scripts/ps/         # Windows run/test/lint scripts
    versions.lock.md    # version pin record
    .github/workflows/  # minimal CI
```

4) Configuration (authoritative)

`core-config.yaml` (already provided) controls:

- **LLM:** `provider: none`, `gate_policy: snippet_only`, `snippet_max_tokens: 220`
- **Budget:** `hard_cap_tokens_per_doc: 1500`, `on_exceed: needs_review`
- **Cache:** SQLite keyed by `(prompt_id, snippet_hash)`
- **OCR:** `enabled: false`
- **Security:** `redact_pii: true` (only applied if LLM enabled)

Dev/QA should never flip these mid-story unless the story explicitly says so.

5) Sharding (critical prep)

Why: agents load **only** what they need.

Commands (BMAD task or CLI):

```
*shard-doc docs/prd.md prd
*shard-doc docs/architecture.md architecture
```

Expected structure:

```
docs/prd/epic-1.md ... docs/architecture/tech_stack.md
                        docs/architecture/coding_standards.md
                        docs/architecture/source_tree.md
                        docs/architecture/api_contracts.md
```

Add the 3 architecture shards + `core-config.yaml` + rulepack files to **dev-load-always**.

6) Agents & Chats (how to drive them)

- **SM chat:** open *SM Story Pack v1* (and Gap Closure Pack) → select the next **Approved** story.
- **Dev chat:** load **that story only** + dev-load-always files → implement.
- **QA chat:** load the same story + repo → verify and mark **Done**.

Do: one story per Dev chat, keep logs, commit often.

Don't: lump multiple stories into one change.

7) Sprint 1 Plan (order & success)

- Order 1. 1.1** Upload & Job Orchestration
2. **1.2** Text Extraction (PDF/DOCX)
 3. **1.3** Evidence Window Builder
 4. **2.1** Rulepack Loader (art28_v1)
 5. **2.2** Detector Runner (verdict + evidence)
 6. **3.1** Findings Table (UI)
 7. **3.2** Report Export (PDF/HTML)
 8. **2.4** Token Ledger & Caps
 9. **4.1** Metrics Wall
 10. **4.2** Coverage Meter
 11. **5.1** Org Settings
 12. **5.2** Minimal Auth & Roles

Sprint success = Upload → Findings → Export works on baseline docs; Metrics shows tiles; Coverage = 8/8.

8) Story Packets (how to use)

Each packet has **Acceptance Criteria**, **Interfaces**, **Tasks**, **Tests**, and **Artifacts**.

- Dev follows tasks exactly; if architecture requires a change, raise to **SM/PO** first.
- QA mirrors **Acceptance Criteria** into a checklist and tests.

The full packets for missing stories (1.3, 2.4, 3.2, 4.1, 4.2, 5.1, 5.2) are in *Gap Closure Pack*.
The rest are in *SM Story Pack v1*.

9) Detection Engine & Rulepack

- **Rulepack:** `apps/api/blackletter_api/rules/art28_v1.yaml` — detectors **(a)–(h)** with anchors, weak cues, and red flags.
- **Lexicon:** `rules/lexicons/weak_language.yaml` — hedges & vagueness.
- **Verdict mapping:** Pass (anchor + no red-flag) / Weak (anchor + hedge) / Missing (no anchor or contradicted) / Needs_review (ambiguous/over budget).
- **Schema validation:** JSON Schema + Pydantic in *Gap Closure Pack*.

Unit tests: minimum 3 **positive** + 3 **hard negative** snippets per detector.

10) Gold Set & Scoring

- **Corpus:** 12–20 contracts (public/synthetic).
- **Labels:** JSONL per finding with verdict + offsets.
- **Scorer CLI:** `tools/score_goldset.py` calculates per-detector Precision/Recall + macro averages; surfaces p95 latency, tokens/doc, %LLM.

Gates: Precision ≥ 0.85 ; Recall ≥ 0.90 ; Explainability $\geq 95\%$; p95 $\leq 60s$; $\leq £0.10/doc$.

11) Frontend UX (what to build)

- **New Analysis:** drag-drop uploader; stepper (Queued→Extracting→Detecting→Reporting→Done); clear error banners.
 - **Findings:** evidence-first table; Verdict chips; Evidence Drawer with snippet + rule id + offsets; filter by verdict; search snippets.
 - **Export:** dialog → PDF/HTML; branding + disclaimers.
 - **Metrics (Admin):** tiles (p95, tokens/doc, %LLM, explainability) + sparklines.
 - **Settings:** LLM provider toggle (default none), OCR, retention; audit changes.
 - **Accessibility:** WCAG AA, keyboard-first, visible focus; color not sole signal.
-

12) Security & Privacy

- LLM **off by default**; snippet-only when enabled; **PII redaction** before any LLM call.
 - Signed URLs, MIME/size checks; retention default **none**.
 - Export includes checksum and timestamp; UI and export show **not legal advice** disclaimer.
-

13) Running Locally (Windows)

- Start both services: `scripts\ps\dev.ps1`
- Run tests: `scripts\ps\test.ps1`

- Lint/format: `scripts\ps\lint.ps1`
- Web: `http://localhost:3000`
- API: `http://localhost:8000`

14) CI & Version Pins

- **CI:** minimal GitHub Actions for backend tests and frontend build/tests (included).
- **Pin policy:** record exact Python and npm trees in `versions.lock.md` on first commit; keep pins stable per sprint.

15) Error Codes & Microcopy (UI/API)

Code	Message	Hint
<code>file_too_large</code>	File too large (max 10MB).	Compress or split the PDF.
<code>mime_unsupported</code>	Unsupported file type.	Upload PDF or DOCX.
<code>extract_failed</code>	Couldn't extract text.	File may be image-only; enable OCR in Settings.
<code>token_cap</code>	Token budget exceeded.	Some checks marked Needs review ; reduce snippet size or raise cap.
<code>report_failed</code>	Report generation failed.	Retry; or download HTML export.

Copy appears in toasts/banners and export footers.

16) Troubleshooting (agentic IDE)

- **Agent stuck:** start a new chat with only the current story + dev-load-always files.
- **Off-spec code:** stop, quote the acceptance criteria back to Dev; if needed, raise to SM to update the packet.
- **Performance drift:** check p95 on Metrics; profile extraction; confirm OCR is off.
- **False greens:** verify rulepack anchors; add hard negatives; widen red-flags.
- **Token cap hits:** confirm provider is **none**; if on, reduce snippet_max or raise cap deliberately.

17) Definition of Ready / Done (DoR / DoD)

- **DoR:** acceptance criteria written; artifacts named; dependencies listed; test data available; UX implications noted.
- **DoD:** unit + integration tests pass; acceptance criteria demo; QA note added; metrics updated; minimal docs/changelog updated.

18) Appendices

A) API Contracts (freeze)

- POST /api/contracts → { job_id, analysis_id, status }
- GET /api/jobs/{job_id} → job status
- GET /api/analyses/{id} → summary + coverage
- GET /api/analyses/{id}/findings → Finding[]
- POST /api/reports/{analysis_id} → { url }

Finding

```
{
  "detector_id": "A28_3_c_security",
  "rule_id": "art28_v1.A28_3_c_security",
  "verdict": "pass|weak|missing|needs_review",
  "snippet": "...technical and organisational measures...",
  "page": 7,
  "start": 1423,
  "end": 1562,
  "rationale": "anchor present; no red-flag",
  "reviewed": false
}
```

B) Commands Cheat Sheet

```
*help                # list BMAD commands
*kb-mode              # load knowledge base
*shard-doc <src> <outdir> # shard markdown by H2
```

C) Story Template (for new items)

```
id: <epic.story>
epic: <n>
title: <short>
status: draft|approved|review|done
acceptance_criteria:
  - <testable condition 1>
  - <testable condition 2>
interfaces:
  - <API/DB/Service contracts>
tasks:
  - <dev tasks>
tests:
  - <unit/integration/UI>
artifacts:
  - <files to create/modify>
```

Final Note

This guide consolidates everything needed to execute **BMAD V4** for Blackletter. Keep agent chats **focused** (one story per tab), and let the documents lead the code. When in doubt, the acceptance criteria are the ground truth.